

# Formula 1 Database

Kishan Patel  
Engineering Sciences (Data Science)  
University at Buffalo  
kpatel39@buffalo.edu

Geethika Bedadhala  
Engineering Sciences (Data Science)  
University at Buffalo  
gbedadhala@buffalo.edu

## I. PROBLEM STATEMENT

To design relational database management system related to Formula1. The system should allow for efficient querying and analysis of data to generate insights and statistics related to F1 races. The overall goal of the database is to provide a comprehensive and accurate representation of F1 race data for analysis and decision-making purposes.

## II. EASE OF USE WHY DO YOU NEED A DATABASE INSTEAD OF AN EXCEL FILE?

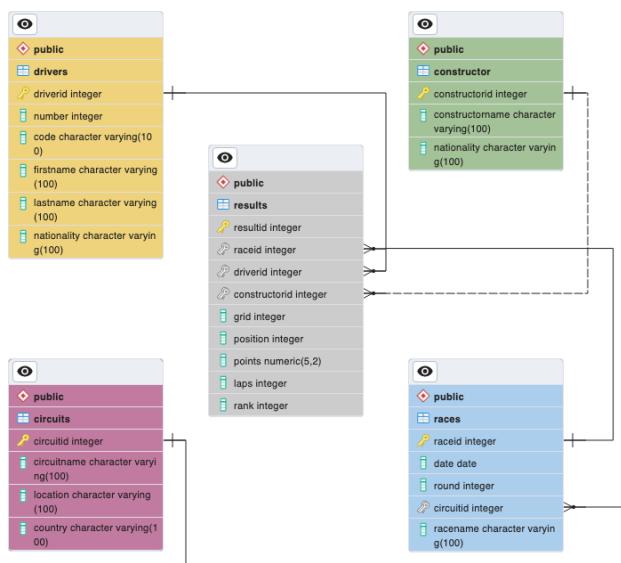
Formula 1 data has large amount of data (more than 1000 - 25000 rows) and multiple excel sheets which are interrelated and excel might be slow and difficult to manage and database can handle large volumes efficiently. Not only that database can provide robust security features than excel and multiple people can access and update the data at the same time easily. Moreover, we can also do effective querying in database than Excel using SQL.

## III. TARGET USERS

The target users of this database system would include F1 fans, race organizers, and media companies. F1 fans can use the database to get information about past and upcoming races, teams, drivers, and circuits. Race organizers can use the database to manage and plan races, track performance metrics, and analyze trends. Media companies can use the database to report on race events, create content, and provide accurate insights to their audiences.

**Database Administer:** Data Analysts who work for Stakeholders.

## IV. DATABASE SCHEMA: ER DIAGRAM



### A. Relations:

- 1) *Constructor* (*ConstructorID*, *constructorname*, *nationality*)

Attributes	Description
<b>ConstructorID</b>	Primary key
<b>ConstructorName</b>	Name of the constructor participating in F1
<b>Nationality</b>	Nationality of the constructor

The above table gives the team details who have participated in Formula 1 from 1950-2023.

- 2) *Drivers* (*DriverID*, *number*, *code*, *firstname*, *lastname*, *nationality*)

Attributes	Description
<b>DriverID</b>	Primary Key
<b>Number</b>	Driver personal car number
<b>Code</b>	Driver name code
<b>FirstName</b>	Driver first name
<b>LastName</b>	Driver last name
<b>Nationality</b>	Driver nationality

The above table gives information of the drivers who have participated in Formula 1 between 1950-2023.

- 3) *Races* (*raceID*, *date*, *round*, *circuitID*, *racename*)

Attributes	Description
<b>RaceID</b>	Primary key
<b>Date</b>	Date of the grand prix race
<b>Round</b>	Round/Race of Formula 1 season
<b>CircuitID</b>	Foreign key from Circuits table
<b>RaceName</b>	Grand Prix name

The above table contains all the race details that have been hosted between 1950-2023.

- 4) *Circuits* (*circuitID*, *circuitname*, *location*, *country*)

Attributes	Description
<b>CircuitID</b>	Primary key

CircuitName	Name of the circuit
Location	Location where circuit is located
Country	Country where the circuit is located

The above table contains the track details where all the races are hosted.

- 5) Results (resultID, raceID, driverID, constructorID, grid, position, points, laps, rank)

Attributes	Description
ResultID	Primary Key
RaceID	Foreign key from Races table
DriverID	Foreign key from Drivers table
ConstructorID	Foreign key from Constructors table
Grid	Grid position the driver started the race in
Position	Position of the driver after the race
Points	Points awarded to the driver according to the position he finished
Laps	Laps finished by the driver after the race
Rank	Rank of the driver in the respective season

The above table contains the results of all the races hosted between 1950-2023 with every detail possible and is connected to all the tables through foreign keys.

## V. SQL QUERIES

- A. Query for finding out Drivers with highest number of wins

```

Query   Query History
1 -- Drivers with race wins.
2
3 SELECT drivers.firstname, drivers.lastname, drivers.nationality
4 FROM drivers
5 WHERE drivers.driverid IN (SELECT drivers.driverid
6                           FROM results
7                           INNER JOIN drivers ON results.driverid = drivers.driverid
8                           WHERE results.position = 1);

```

Data Output Messages Notifications

firstname	lastname	nationality
Lewis	Hamilton	British
Nico	Rosberg	German
Fernando	Alonso	Spanish
Heikki	Kovalainen	Finnish
Kimi	Räikkönen	Finnish
Robert	Kubica	Polish
Felipe	Massa	Brazilian
David	Coulthard	British
Jarno	Trulli	Italian
Mark	Webber	Australian
Jenson	Button	British
Sebastian	Vettel	German
Giancarlo	Fisichella	Italian
Rubens	Barrichello	Brazilian
Ralf	Schumacher	German
Michael	Schumacher	German
Juan	Pablo Montoya	Colombian
Jacques	Villeneuve	Canadian
Olivier	Panis	French
Heinz-Harald	Frentzen	German
Jean	Alesi	French
Eddie	Irvine	British
Mika	Häkkinen	Finnish
Johnny	Herbert	British
Damon	Hill	British
Gerhard	Berger	Austrian

Total rows: 113 of 113 | Query complete 00:00:00.106

- B. query for finding out Drivers with race wins(only 5 rows displayed)

```

Query   Query History
1 -- Drivers with highest number of race wins.
2
3 SELECT drivers.firstname, drivers.lastname, drivers.nationality, COUNT(*) AS RaceWins
4 FROM results
5 INNER JOIN drivers ON results.driverid = drivers.driverid
6 WHERE results.position = 1
7 GROUP BY drivers.driverid
8 ORDER BY COUNT(*) DESC
9 LIMIT 10;

```

Data Output Messages Notifications

firstname	lastname	nationality	racewins
Lewis	Hamilton	British	103
Michael	Schumacher	German	91
Sebastian	Vettel	German	53
Alain	Prost	French	51
Ayrton	Senna	Brazilian	41
Max	Verstappen	Dutch	35
Fernando	Alonso	Spanish	32
Nigel	Mansell	British	31
Jackie	Stewart	British	27
Jim	Clark	British	25

### C. query for finding out constructors with race wins(only 5 rows displayed)

Query History	
1 -- Constructors with race wins.	
2 SELECT constructor.constructorname as Constructor, constructor.Nationality	
3 FROM constructor	
4 WHERE constructor.constructorid IN (SELECT constructor.constructorid	
5 FROM results	
6 INNER JOIN constructor on results.constructorid = constructor.constructorid	
7 WHERE results.position = 1);	
Data Output	Messages Notifications

	constructor	nationality
1	McLaren	British
2	BMW Sauber	German
3	Williams	British
4	Renault	French
5	Toro Rosso	Italian

### D. query for finding out Fernando Alonsos race wins

Query History	
1 -- Fernando Alonso's race wins.	
2 SELECT races.racename, circuits.circuitname, races.date, results.grid as StartPosition, constructor.constructorname	
3 FROM results	
4 INNER JOIN drivers on results.driverid = drivers.driverid	
5 INNER JOIN constructor on results.constructorid = constructor.constructorid	
6 INNER JOIN races on results.raceid = races.raceid	
7 INNER JOIN circuits on races.circuitid = circuits.circuitid	
8 WHERE results.position = 1 and drivers.firstname = 'Fernando' and drivers.lastname = 'Alonso'	
9 ORDER BY races.date;	
Data Output	Messages Notifications

	racename	circuitname	date	startposition	constructorname
1	Hungarian Grand Prix	Hungaroring	2003-08-24	1	Renault
2	Malaysian Grand Prix	Sepang International Circuit	2006-03-20	1	Renault
3	Bahrain Grand Prix	Bahrain International Circuit	2005-04-03	1	Renault
4	San Marino Grand Prix	Autodromo Enzo e Dino Ferrari	2005-04-24	2	Renault
5	European Grand Prix	Nürburgring	2005-05-29	6	Renault

## VI BCNF/3NF CHECK

For a relation to be in BCNF from, it should satisfy the following conditions:

- 1) No nontrivial functional dependencies of non-prime attributes on any candidate key exist.
- 2) Every nontrivial functional dependency is a dependency on a candidate key.

Let us check whether the above relations are in BCNF form or not:

### A. Constructor (ConstructorID, constructorname, nationality)

The above relation is in BCNF form with one candidate key i.e., ConstructorID.

FDs of the relation are:

- ConstructorID → ConstructorName
- ConstructorID → Nationality

Both FDs are trivial because the left-hand side of each FD is the primary key.

### B. Drivers (DriverID, number, code, firstname, lastname, nationality)

The above relation is in BCNF form with one candidate key i.e., DriverID.

FDs of the relation are:

- DriverID → number
- DriverID → code
- DriverID → firstname
- DriverID → lastname
- DriverID → nationality

Above FDs are trivial because the left-hand side of each FD is the primary key.

### C. Circuits (circuitID, circuitname, location, country)

The above relation is in BCNF form with one candidate key i.e., CircuitID.

FDs of the relation are:

- CircuitID → circuitname
- CircuitID → location
- CircuitID → country

Above FDs are trivial because the left-hand side of each FD is the primary key.

### D. Races (raceID, date, round, circuitID, racename)

The above relation is not in BCNF form. To be in BCNF, every determinant of the relation must be a candidate key. In the given relation, the candidate key is (RaceID), and the only non-prime attribute is RaceName. The determinant (CircuitID) is not a candidate key and is a non-prime attribute.

To convert the relation into BCNF, we need to decompose it into two relations:

- Races (RaceID, Date, Round, RaceName)
- RaceCircuits (RaceID, CircuitID)

But the above relation is in 3NF form as it satisfies all the conditions of 3NF. We would like to keep the above relation as decomposing it into BCNF would result in two tables, which may be more complex than necessary for some applications. To keep a simpler design for the database, keeping it in 3NF form seems like a better option.

### E. Results (resultID, raceID, driverID, constructorID, grid, position, points, laps, rank)

The above relation is in BCNF form with one candidate key i.e., ResultID.

FDs of the relation are:

- ResultID → RaceID, DriverID, ConstructorID, grid, position, points, laps, rank

Above FD is trivial because the left-hand side of FD is a primary key.

Overall schema satisfies the BCNF requirement.

## VII FINALIZING THE E/R DIAGRAM

Since our data is already in BCNF form there will be no change in ER Diagram.

## VIII PROBLEMS WHILE USING LARGE DATABASE AND USING INDEXING TO OVERCOME THIS.

Indexing is required in our database to improve query performance. The performance of queries that involve the indexed attributes may improve significantly. This is because the index creates a separate data structure that

allows the database management system to find records more quickly and efficiently.

However, we should not create too many indexes as it can also have a negative impact on query performance, as it can slow down data modification operations on the database. We have applied indexing in following ways for all relations:

1) ‘Results’ relation has 3 foreign keys. Hence to improve query performance when joining tables, we have created indexes for all three foreign keys i.e., raceid, driverid, constructorid.

2) Similarly, for ‘Races’ relation, we have applied indexing to circuitid column.

3) There might be multiple queries in the future, that frequently filter or sort the ‘Drivers’ or ‘Constructors’ tables by the ‘Nationality’ attribute, hence we created indexing for nationality attribute for both the relations.

#### A. Selecting type of indexing:

Here we used Bitmap indexing. Since our data is large creating dense indexing does not make sense as it’s going to take more time to execute the query as well. Sparse data is also the same reason why we didn’t use it. The B-tree index in PostgreSQL is effective for range queries, equality queries, and ordered retrieval of data. It organizes the data in a hierarchical structure, allowing efficient searching, insertion, deletion, and updates. The B-tree index is suitable for a wide range of data types, including integers, floating-point numbers, strings, and timestamps. In F1 race DataBase we often query that uses range and ordered retrieval. So B-tree indexing is best suited for our query. Hash indexing can be a bit redundant as we often do not execute exact match queries.

#### B. Executing 3 complex queries and comparing cost before and after indexing

##### 1) query 1

###### a) before

```
Query - Query History
1 EXPLAIN SELECT Drivers.FirstName, Drivers.LastName, Races.Date, Circuits.Location, MAX(Results.points) as MaxPoints
2 FROM Results
3 JOIN Races ON Races.RaceID = Results.RaceID
4 JOIN Drivers ON Drivers.DriverID = Results.DriverID
5 JOIN Circuits ON Circuits.CircuitID = Races.CircuitID
6 GROUP BY Results.RaceID, Drivers.FirstName, Drivers.LastName, Races.Date, Circuits.Location
7 LIMIT 10

Data Output Messages Notifications
QUERY PLAN
text
1 Limit (cost=991.56 553 rows=10 width=62)
2 > GroupAggregate (cost=991.56 431903.28 rows=25840 width=62)
3 Group Key: results.raceid, drivers.firstname, drivers.lastname, races.date, circuits.location
4 > Incremental Sort (cost=991.56 .43125/28 rows=25840 width=33)
5 Sort Key: results.raceid, drivers.firstname, drivers.lastname, races.date, circuits.location
6 Presorted Key: results.raceid
7 > Nested Loop (cost=0.56 .429884.51 rows=25840 width=33)
8 > Nested Loop (cost=0.28 .429887.02 rows=25840 width=24)
9 Join Filter: (races.raceid = results.raceid)
10 > Nested Loop (cost=0.28 .1308.82 rows=1102 width=17)
11 Join Filter: (races.circuitid = circuits.circuitid)
12 > Index Scan using races_pkey on races (cost=0.28 .49.81 rows=1102 width=12)
13 > Materialize (cost=0.00 .2.16 rows=77 width=13)
14 > Seq Scan on circuits (cost=0.00 .0.35 rows=1 width=13)
15 > Materialize (cost=0.00 .607.60 rows=25840 width=11)
16 > Seq Scan on results (cost=0.00 .478.40 rows=25840 width=11)
17 > Memzome (cost=0.29 .0.30 rows=1 width=17)
18 Cache Key: results.raceid
20 > Index Scan using drivers_pkey on drivers (cost=0.28 .0.29 rows=1 width=17)
Index Cond: (driverid = results.driverid)
```

##### b) after

```
Query - Query History
1 EXPLAIN SELECT Drivers.FirstName, Drivers.LastName, Races.Date, Circuits.Location, MAX(Results.points) as MaxPoints
2 FROM Results
3 JOIN Races ON Races.RaceID = Results.RaceID
4 JOIN Drivers ON Drivers.DriverID = Results.DriverID
5 JOIN Circuits ON Circuits.CircuitID = Races.CircuitID
6 GROUP BY Results.RaceID, Drivers.FirstName, Drivers.LastName, Races.Date, Circuits.Location
7 LIMIT 10

Data Output Messages Notifications
QUERY PLAN
text
1 Limit (cost=0.79 .537 rows=10 width=62)
2 > GroupAggregate (cost=79.4093.38 rows=25840 width=62)
3 Group Key: results.raceid, drivers.firstname, drivers.lastname, races.date, circuits.location
4 > Incremental Sort (cost=79.3447.38 rows=25840 width=33)
5 Sort Key: results.raceid, drivers.firstname, drivers.lastname, races.date, circuits.location
6 Presorted Key: results.raceid
7 > Nested Loop (cost=1.20 .2074.61 rows=25840 width=23)
8 > Merge Join (cost=0.72 .1177.12 rows=25840 width=24)
9 Index Cond: (races.raceid = results.raceid)
10 > Nested Loop (cost=0.48 .99.44 rows=1102 width=17)
11 > Index Scan using races_pkey on races (cost=0.28 .49.81 rows=1102 width=12)
12 > Memzome (cost=0.15 .0.18 rows=1 width=13)
13 Cache Key: races.circuitid
14 Cache Mode: logical
15 > Index Scan using circuits_pkey on circuits (cost=0.14 .0.17 rows=1 width=13)
16 Index Cond: (circuitid = races.circuitid)
17 > Index Scan using fk_raced_on results (cost=0.29 .762.78 rows=25840 width=11)
18 > Memzome (cost=0.29 .0.30 rows=1 width=17)
19 Cache Key: results.driverid
20 Cache Mode: logical
21 > Index Scan using drivers_pkey on drivers (cost=0.28 .0.29 rows=1 width=17)
22 Index Cond: (driverid = results.driverid)
```

Reason: We have indexed Race ID and Driver ID and index scanning shouldn’t take much time to retrieve. If it’s not present, then every row of the data has to be checked before grouping.

#### 2). Query 2

###### a) before

```
Query - Query History
1 EXPLAIN SELECT races.racename, circuits.circutname, races.date, results.grid as StartPosition, constructor.constructorname
2 FROM results
3 INNER JOIN drivers on results.driverid = drivers.driverid
4 INNER JOIN constructor on results.constructorid = constructor.constructorid
5 INNER JOIN races on results.raceid = races.raceid
6 INNER JOIN circuits on races.circuitid = circuits.circuitid
7 WHERE results.position = 1 and drivers.firstname = 'Fernando' and drivers.lastname = 'Alonso'
8 ORDER BY races.date

Data Output Messages Notifications
QUERY PLAN
text
1 Sort (cost=564.43 .564.43 rows=1 width=57)
2 Sort Key: races.date
3 > Nested Loop (cost=20.43 .366.42 rows=1 width=57)
4 > Nested Loop (cost=20.28 .366.25 rows=1 width=49)
5 > Nested Loop (cost=20.01 .365.90 rows=1 width=17)
6 > Hash Join (cost=19.87 .365.72 rows=1 width=17)
7 Hash Cond: (results.driverid = drivers.driverid)
8 > Seq Scan on results (cost=0.00 .543.00 rows=1082 width=16)
9 Filter: ('position' = 1)
10 > Hash (cost=19.85 .19.85 rows=1 width=4)
11 > Seq Scan on drivers (cost=0.00 .19.85 rows=1 width=4)
12 Filter: ((firstname) = 'Fernando' AND (lastname) = 'Alonso')
13 > Index Scan using constructor_pkey on constructor (cost=0.14 .0.18 rows=1 width=13)
14 Index Cond: (constructorid = results.constructorid)
15 > Index Scan using races_pkey on races (cost=0.28 .0.35 rows=1 width=31)
16 Index Cond: (raced = results.raceid)
17 > Index Scan using circuits_pkey on circuits (cost=0.14 .0.17 rows=1 width=25)
18 Index Cond: (circuitid = races.circuitid)
```

###### b) after

```
Query - Query History
1 EXPLAIN SELECT races.racename, circuits.circutname, races.date, results.grid as StartPosition, constructor.constructorname
2 FROM results
3 INNER JOIN drivers on results.driverid = drivers.driverid
4 INNER JOIN constructor on results.constructorid = constructor.constructorid
5 INNER JOIN races on results.raceid = races.raceid
6 INNER JOIN circuits on races.circuitid = circuits.circuitid
7 WHERE results.position = 1 and drivers.firstname = 'Fernando' and drivers.lastname = 'Alonso'
8 ORDER BY races.date

Data Output Messages Notifications
QUERY PLAN
text
1 Sort (cost=93.09 .93.09 rows=1 width=57)
2 Sort Key: races.date
3 > Nested Loop (cost=21.85 .30.21 rows=1 width=57)
4 > Nested Loop (cost=21.71 .30.21 rows=1 width=49)
5 > Nested Loop (cost=21.43 .31.25 rows=1 width=17)
6 > Nested Loop (cost=21.29 .32.26 rows=1 width=12)
7 > Seq Scan on drivers (cost=0.00 .19.85 rows=1 width=4)
8 Filter: ((firstname) = 'Fernando' AND (lastname) = 'Alonso')
9 > Index Scan using drivers_pkey on drivers (cost=0.29 .72.51 rows=1 width=16)
10 Index Cond: (driverid = results.driverid)
11 Filter: ('position' = 1)
12 > Index Scan using constructor_pkey on constructor (cost=0.14 .0.18 rows=1 width=13)
13 Index Cond: (constructorid = results.constructorid)
14 > Index Scan using races_pkey on races (cost=0.28 .0.35 rows=1 width=31)
15 Index Cond: (raced = results.raceid)
16 > Index Scan using circuits_pkey on circuits (cost=0.14 .0.17 rows=1 width=25)
17 Index Cond: (circuitid = races.circuitid)
```

Reason: here order by took time to execute but that significantly reduced after indexing

## IX SQL QUERIES

Query    Query History

```

1  -- Query to show 5th round of Formula 1 championship season every year
2  SELECT Races.RaceName, Circuits.Location, Circuits.Country
3  FROM Races
4  INNER JOIN Circuits ON Races.CircuitID = Circuits.CircuitID
5  WHERE Races.Round = 5;

```

Data Output    Messages    Explain    X    Notifications

	raceName	location	country
1	Spanish Grand Prix	Montmeló	Spain
2	Turkish Grand Prix	Istanbul	Turkey
3	Monaco Grand Prix	Monte-Carlo	Monaco
4	European Grand Prix	Nürburg	Germany
5	Spanish Grand Prix	Montmeló	Spain

Query    Query History

```

1  -- Query showing circuit names and the number of times they have hosted a race
2  SELECT Circuits.CircuitName, COUNT(Races.RaceID) as NumberOfRaces
3  FROM Circuits
4  JOIN Races ON Races.CircuitID = Circuits.CircuitID
5  GROUP BY Circuits.CircuitID
6  ORDER BY NumberOfRaces DESC

```

Data Output    Messages    Explain    X    Notifications

	circuitName	numberofraces
1	Autodromo Nazionale di Monza	73
2	Circuit de Monaco	69
3	Silverstone Circuit	58
4	Circuit de Spa-Francorchamps	56
5	Circuit Gilles Villeneuve	42

Query    Query History

```

1  -- Query listing all the constructors with their respective number of race wins in F1.
2  SELECT Constructor.ConstructorName, COUNT(*) AS wins
3  FROM Constructor
4  INNER JOIN Results ON Constructor.ConstructorID = Results.ConstructorID
5  WHERE Results.position = 1
6  GROUP BY Constructor.ConstructorName
7  ORDER BY wins DESC;

```

Data Output    Messages    Explain    X    Notifications

	constructorName	wins
1	Ferrari	243
2	McLaren	179
3	Mercedes	125
4	Williams	114
5	Red Bull	92

Query    Query History

```
1 -- Query showing results of US GP 2022.
2 SELECT Drivers.FirstName, Drivers.LastName, Constructor.ConstructorName, Results.position
3 FROM Results
4 JOIN Races ON Races.RaceID = Results.RaceID
5 JOIN Drivers ON Drivers.DriverID = Results.DriverID
6 JOIN Constructor ON Constructor.ConstructorID = Results.ConstructorID
7 WHERE Races.Date = '2022-10-23';
```

Data Output    Messages    Explain    Notifications



	firstname character varying (100)	lastname character varying (100)	constructorname character varying (100)	position integer
1	Max	Verstappen	Red Bull	1
2	Lewis	Hamilton	Mercedes	2
3	Charles	Leclerc	Ferrari	3
4	Sergio	Pérez	Red Bull	4
5	George	Russell	Mercedes	5

Query    Query History

```
1 -- PROCEDURE: public.deletedriver(integer)
2
3 DROP PROCEDURE IF EXISTS public.deletedriver(integer);
4
5 CREATE OR REPLACE PROCEDURE public.deletedriver(
6     IN did integer)
7 LANGUAGE 'plpgsql'
8 AS $BODY$
9 BEGIN
10    DELETE FROM Results WHERE DriverID = dID;
11    DELETE FROM Drivers WHERE DriverID = dID;
12 END;
13 $BODY$;
14 ALTER PROCEDURE public.deletedriver(integer)
15     OWNER TO postgres;
16
17 COMMENT ON PROCEDURE public.deletedriver(integer)
18     IS 'Procedure to delete a specific driver and all of their race results.';
19
20 CALL public.deletedriver(
21     14
22 )
```

Data Output    Messages    Notifications

CALL

Query returned successfully in 42 msec.

Query    Query History

```

1 -- PROCEDURE: public.update_nationality(integer, text, text)
2
3 -- DROP PROCEDURE IF EXISTS public.update_nationality(integer, text, text);
4
5 CREATE OR REPLACE PROCEDURE public.update_nationality(
6     IN id integer,
7     IN table_name text,
8     IN new_nationality text)
9 LANGUAGE 'plpgsql'
10 AS $BODY$
11 BEGIN
12 IF table_name = 'drivers' THEN
13     UPDATE drivers SET nationality = new_nationality WHERE driverid = id;
14 ELSIF table_name = 'constructor' THEN
15     UPDATE constructor SET nationality = new_nationality WHERE constructorid = id;
16 END IF;
17 END;
$BODY$;
19 ALTER PROCEDURE public.update_nationality(integer, text, text)
20     OWNER TO postgres;
21
22 COMMENT ON PROCEDURE public.update_nationality(integer, text, text)
23     IS 'Procedure to update the nationality of a specific constructor or driver.';
24
25
26 CALL public.update_nationality(
27     1,
28     'drivers',
29     'German'
30 )

```

Data Output    Messages    Notifications

CALL

Query returned successfully in 53 msec.

Query    Query History

```

1 -- FUNCTION: public.get_most_raced_circuit()
2
3 -- DROP FUNCTION IF EXISTS public.get_most_raced_circuit();
4
5 CREATE OR REPLACE FUNCTION public.get_most_raced_circuit(
6     )
7     RETURNS text
8     LANGUAGE 'plpgsql'
9     COST 100
10    VOLATILE PARALLEL UNSAFE
11   AS $BODY$
12  BEGIN
13      RETURN (
14          SELECT circuits.circuitname
15          FROM circuits
16          JOIN races ON circuits.circuitid = races.circuitid
17          GROUP BY circuits.circuitid, circuits.circuitname
18          ORDER BY COUNT(races.raceid) DESC
19          LIMIT 1
20      );
21  END;
$BODY$;
23
24 ALTER FUNCTION public.get_most_raced_circuit()
25     OWNER TO postgres;
26
27 COMMENT ON FUNCTION public.get_most_raced_circuit()
28     IS 'A function to return the circuit with the most races: This function would
29     join the "Races" and "Circuits" tables and group the races by circuit, counting
30     the number of races for each circuit. The function would return the circuit with
31     the highest number of races.';
32
33 SELECT public.get_most_raced_circuit();

```

Data Output    Messages    Notifications

	get_most_raced_circuit	text	🔒
1	Autodromo Nazionale di Monza		

Query    Query History

```

1 -- FUNCTION: public.get_driver_season_points(integer, integer)
2
3 -- DROP FUNCTION IF EXISTS public.get_driver_season_points(integer, integer);
4
5 CREATE OR REPLACE FUNCTION public.get_driver_season_points(
6   driver_id integer,
7   season_year integer)
8   RETURNS numeric
9   LANGUAGE 'plpgsql'
10  COST 100
11 VOLATILE PARALLEL UNSAFE
12 AS $BODY$
13 $BEGIN
14   RETURN (
15     SELECT sum(results.points)
16       FROM results
17      JOIN races ON results.raceid = races.raceid
18     WHERE results.driverid = driver_id AND extract(YEAR FROM races.date) = season_year
19   );
20 $END;
21 $BODY$;
22
23 ALTER FUNCTION public.get_driver_season_points(integer, integer)
24   OWNER TO postgres;
25
26 COMMENT ON FUNCTION public.get_driver_season_points(integer, integer)
27   IS 'A function to calculate the total points earned by a driver in a given season:
28   This function would take the DriverID and a season year as input parameters and would
29   calculate the total points earned by the driver in all races for that season.
30   The function would need to join the "Results" and "Races" tables to filter results by
31   the season year and driver ID, and then sum the "points" attribute for all results matching the criteria.';
32
33
34 SELECT public.get_driver_season_points(
35   1, 2018
36 )
37

```

Data Output    Messages    Notifications

	get_driver_season_points	numeric
1		408.00

Query    Query History

```

1 -- Trigger to set "position" to "DNF" when a driver has not completed all laps.
2
3 DROP TRIGGER IF EXISTS set_position_dnf_trigger ON public.results;
4
5 CREATE OR REPLACE FUNCTION set_position_dnf()
6   RETURNS trigger AS
7 $BODY$
8 $BEGIN
9   IF NEW.laps < (SELECT MAX(laps) FROM Results WHERE Results.RaceID = NEW.RaceID) THEN
10     NEW.position = 'DNF';
11   END IF;
12   RETURN NEW;
13 END;
14 $BODY$;
15 LANGUAGE plpgsql;
16
17 CREATE TRIGGER set_position_dnf_trigger
18 BEFORE INSERT ON Results
19 FOR EACH ROW
20 EXECUTE FUNCTION set_position_dnf();

```

Query    Query History

```

1 SELECT * FROM public.results
2 ORDER BY resultid ASC

```

Data Output    Messages    Notifications

resultid [PK] integer	raceid integer	driverid integer	constructorid integer	grid integer	position character varying (20)	points numeric (5,2)	laps integer	rank integer
1	1	18	1	1	1	10.00	58	2
2	2	18	2	2	5 2	8.00	58	3
3	3	18	3	3	7 3	6.00	58	5
4	4	18	4	4	11 4	5.00	58	7
5	5	18	5	1	3 5	4.00	58	1
6	6	18	6	3	13 DNF	3.00	57	14
7	7	18	7	5	17 DNF	2.00	55	12
8	8	18	8	6	15 DNF	1.00	53	4
9	9	18	9	2	2 DNF	0.00	47	9
10	10	18	10	7	18 DNF	0.00	43	13
11	11	18	11	8	19 DNF	0.00	32	15
12	12	18	12	4	20 DNF	0.00	30	16
13	13	18	13	6	4 DNF	0.00	29	6
14	14	18	14	9	8 DNF	0.00	25	11
15	15	18	15	7	6 DNF	0.00	19	10
16	16	18	16	10	22 DNF	0.00	8	17
17	17	18	17	9	14 DNF	0.00	0	[null]
18	18	18	18	11	12 DNF	0.00	0	[null]
19	19	18	19	8	21 DNF	0.00	0	[null]
20	20	18	20	5	9 DNF	0.00	0	[null]
21	21	18	21	10	16 DNF	0.00	0	[null]
22	22	18	22	11	10 22	0.00	58	8
23	23	19	8	6	2 1	10.00	56	2
24	24	19	9	2	4 2	8.00	56	6
25	25	19	5	1	8 3	6.00	56	7
26	26	19	15	7	3 4	5.00	56	8
27	27	19	1	1	9 5	4.00	56	3
28	28	19	2	2	5 6	3.00	56	1
29	29	19	17	9	6 7	2.00	56	12
30	30	19	4	4	7 8	1.00	56	10

Total rows: 1000 of 25840    Query complete 00:00:00.524

Query    Query History

---

```
1 -- Some Insert, Delete and Update Queries.
2 INSERT INTO Constructor (ConstructorID, ConstructorName, Nationality)
3 VALUES (212, 'Buffalo', 'USA');
4
5 UPDATE Circuits
6 SET Country = 'Spain'
7 WHERE CircuitName = 'Circuit de Barcelona-Catalunya';
8
9 INSERT INTO Races (RaceID, Date, Round, CircuitID, RaceName)
10 VALUES (2106, '2023-06-18', 9, 2, 'French Grand Prix');
11
12 UPDATE Results
13 SET points = 12
14 WHERE ResultID = 1;
15
16 DELETE FROM RESULTS
17 WHERE ConstructorID = 3;
```

---

Data Output    Messages    Notifications

DELETE 140

Query returned successfully in 31 msec.

*Contribution: Kishan Patel 50%, Geethika Bedadhala 50%. We both were in sync throughout the project.*